

TGS Saltbody Detection (Kaggle)

Predict saltbody on multiple 2D seismic images by U-Net

Import Libraries

```
In [1]: import time
import os
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import imageio
import copy
import tensorflow as tf
from keras import backend as K
from keras import Input, optimizers
from keras.models import Model
from keras.losses import binary_crossentropy
from keras.utils.vis_utils import plot_model
from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau
from keras.layers import Conv2D, MaxPooling2D, UpSampling2D, Dropout, concatenate
from keras_tqdm import TQDMNotebookCallback
from sklearn.model_selection import train_test_split
from skimage.transform import resize
```

Using TensorFlow backend.

Data Loading

```
In [2]: path_data = os.getcwd() + '/dataset'
path_train_img = path_data + '/train/images/'
path_train_msk = path_data + '/train/masks/'
path_test_img = path_data + '/test/'
train = pd.read_csv(path_data + '/train.csv')
sample_submit = pd.read_csv(path_data + '/sample_submission.csv')

size_img = 101
size_img_target = 128
num_epoch = 50
num_batch = 32
```

```
In [3]: def load_images(size_img, path_img, path_msk=[]):
    ''' Load image files into memory '''
    flatten_msk = lambda x: list(map(int, np.reshape(x, np.size(x, 0) * np.size(x, 1))))
    num_imgs = len([name for name in os.listdir(path_img) if name.endswith('.png')])
    i = 0
    ids = []
    data_img = np.zeros([num_imgs, size_img, size_img])
    data_msk = np.zeros([num_imgs, size_img, size_img])

    for file in os.listdir(path_img):
        if file.endswith('.png'):
            id_tmp = file[:-4]
            ids.append(id_tmp)
            data_img[i, :, :] = np.array(list(zip(*imageio.imread(path_img + file)[: , :, 0] / 255)))
            if path_msk:
                msk = flatten_msk(list(zip(*imageio.imread(path_msk + file))))
                msk = [1 if i > 0 else 0 for i in msk]
                data_msk[i, :, :] = np.array(msk).reshape([size_img, size_img])
            i += 1
    return ids, data_img, data_msk
```

```
In [4]: ids_train, data_img, data_msk = load_images(size_img, path_train_img, path_train_msk)
ids_test, test_img, _ = load_images(size_img, path_test_img)
```

Data Preprocessing (Upsampling)

```
In [5]: def upsample(img):
        ''' Upsample raw images from 101 x 101 to 128 x 128 '''
        if size_img == size_img_target:
            return img
        return resize(img, (size_img_target, size_img_target), mode='constant', preserve_range=True)

        util_upsample = lambda data: np.array(list(map(upsample,data))).reshape(-1,size_img_target,size_img_target,1)
```

```
In [6]: data_img_up = util_upsample(data_img)
        data_msk_up = util_upsample(data_msk)
        test_img_up = util_upsample(test_img)
        x_train = np.append(data_img_up, [np.fliplr(x) for x in data_img_up], axis=0)
        y_train = np.append(data_msk_up, [np.fliplr(x) for x in data_msk_up], axis=0)
        im_train,im_val,lb_train,lb_val=train_test_split(x_train,y_train,test_size=0.2,random_state=42)
```

Model Metrics

```
In [7]: def dice_loss(y_true, y_pred):
        ''' Dice Loss function '''
        smooth = 1.
        y_true_f = K.flatten(y_true)
        y_pred_f = K.flatten(y_pred)
        intersection = y_true_f * y_pred_f
        score = (2. * K.sum(intersection) + smooth) / (K.sum(y_true_f) + K.sum(y_pred_f) + smooth)
        return 1. - score
```

```
In [8]: def bce_dice_loss(y_true, y_pred):
        ''' Binary Cross Entropy (BCE) Dice Loss function '''
        return binary_crossentropy(y_true, y_pred) + dice_loss(y_true, y_pred)
```

Model (U-Net)

```
In [9]: def Conv2d_with_prm(model_in, num_filters):
        ''' Wrapper for Conv2D function '''
        conv = Conv2D(filters=num_filters,
                      kernel_size=(3,3),
                      activation='relu',
                      padding='same',
                      kernel_initializer='he_normal')(model_in)

        return conv
```

```
In [10]: def conv2d_down(model_in, num_filters, pooling=False, dropout=False):
        ''' Downward part for U-Net '''
        drop = []
        pool = []
        conv = Conv2d_with_prm(model_in, num_filters)
        conv = Conv2d_with_prm(conv, num_filters)
        if dropout:
            drop = Dropout(rate=0.5)(conv)
        if pooling:
            pool = MaxPooling2D(pool_size=(2, 2))(conv)
        return conv, drop, pool
```

```
In [11]: def conv2d_up(model_in, model_merge, num_filters):
        ''' Upward part for U-Net '''
        up = Conv2d_with_prm(UpSampling2D(size=(2,2))(model_in) , num_filters)
        merge = concatenate([model_merge, up], axis=3)
        conv = Conv2d_with_prm(merge, num_filters)
        conv = Conv2d_with_prm(conv, num_filters)
        return conv, up, merge
```

```
In [12]: def visualize_model(self, path_savefig='model_plot.png'):
        ''' Summarize and visualize created U-Net model '''
        self.model.summary()
        plot_model(self.model, to_file=path_savefig, show_shapes=True, show_layer_names=True)
```

```
In [13]: def unet(input_size,lr=1e-4):
''' Compile U-Net Model '''
inputs = Input(input_size)
conv1, _ = conv2d_down(inputs, 2**4, pooling=True)
conv2, _ = conv2d_down(pool1, 2**5, pooling=True)
conv3, _ = conv2d_down(pool2, 2**6, pooling=True)
_, drop4, pool4 = conv2d_down(pool3, 2**7, pooling=True, dropout=True)
_, drop5, _ = conv2d_down(pool4, 2**8, pooling=False, dropout=True)
conv6, _ = conv2d_up(drop5, drop4, 2**7)
conv7, _ = conv2d_up(conv6, conv3, 2**6)
conv8, _ = conv2d_up(conv7, conv2, 2**5)
conv9, _ = conv2d_up(conv8, conv1, 2**4)
conv9 = Conv2D(2, 3, activation='relu', padding='same', kernel_initializer='he_normal')(conv9)
conv10 = Conv2D(1, 1, activation='sigmoid')(conv9)
model = Model(inputs=inputs, outputs=conv10)
model.compile(optimizer=optimizers.Adam(lr),loss=bce_dice_loss, metrics=['acc'])
return model
```

```
In [14]: tf.logging.set_verbosity(tf.logging.ERROR)
model = unet((size_img_target,size_img_target,1),lr=1e-3)
early_stopping = EarlyStopping(patience=10, verbose=0)
model_checkpoint = ModelCheckpoint('./keras.model', save_best_only=True, verbose=0)
reduce_lr = ReduceLRonPlateau(factor=0.1, patience=4, min_lr=0.00001, verbose=1)

tic = time.time()
history = model.fit(im_train,lb_train,batch_size=num_batch,epochs=num_epoch,
                    validation_data=(im_val,lb_val),verbose=0,shuffle=True,
                    callbacks=[early_stopping, model_checkpoint, reduce_lr, TQDMNotebookCallback()])
toc = time.time()
```

Training 78% 39/50 [09:06<02:32, 13.86s/it]

Epoch 00028: ReduceLRonPlateau reducing learning rate to 0.00010000000474974513.

Epoch 00033: ReduceLRonPlateau reducing learning rate to 1.0000000474974514e-05.

Epoch 00037: ReduceLRonPlateau reducing learning rate to 1e-05.

```
In [15]: print('Comutation time: ' + '{:1.1f}'.format((toc - tic)/60) + ' min')
```

Comutation time: 9.2 min

Training Metrics for every epoch


```
In [16]: def plot_acc_epoch(acc,val_acc,loss,val_loss):
''' Plot accuracy and loss for training and validation data '''

# Accuracy
epochs = range(1,len(acc) + 1)
plt.subplots(1, 2, figsize=(12, 5))
ax1 = plt.subplot(121)
plt.plot(epochs,acc,'b',linewidth=1.5,label='Training acc')
plt.plot(epochs,val_acc,'r',linewidth=1.5,label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
ax1.set_ylim([0.7,1])
plt.legend(loc='lower right')

# Loss
ax2 = plt.subplot(122)
plt.plot(epochs,loss,'b',linewidth=1.5,label='Training loss')
plt.plot(epochs,val_loss,'r',linewidth=1.5,label='Validation loss')
plt.title('Training and validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.show()
```

```
In [17]: acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
```

```
In [18]: plot_acc_epoch(acc, val_acc, loss, val_loss)
```



Visualize some of the trained weights

```
In [19]: def deprocess_image(x):
''' Utility function to convert a tensor into a valid image '''
x -= x.mean()
x /= (x.std() + 1e-5)
x *= 0.1
x += 0.5
x = np.clip(x, 0, 1)
x *= 255
x = np.clip(x, 0, 255).astype('uint8')
return x
```

```
In [20]: def generate_pattern(model, layer_output, filter_index, size):
''' Function to generate filter visualizations '''
loss = K.mean(layer_output[:, :, :, filter_index])
grads = K.gradients(loss, model.input)[0]
grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)
iterate = K.function([model.input], [loss, grads])
input_img_data = np.random.random((1, size, size, 1)) * 20 + 128


step = 1.
for i in range(40):
    loss_value, grads_value = iterate([input_img_data])
    input_img_data += grads_value * step

img = input_img_data[0]
return np.reshape(deprocess_image(img), [size, size])
```


```
In [21]: def plt_filter_patterns(model, layer_name):
''' Plot filter patterns '''
layer_output = model.get_layer(layer_name).output
size = layer_output.get_shape().as_list()[1]
margin = int(5 - np.log2(128/size))
results = np.zeros((2*size + 1*margin, 8*size + 7*margin))
for i in range(2):
    for j in range(8):
        filter_img = generate_pattern(model, layer_output, j + (i * 8), size)
        horizontal_start = i * size + i * margin
        horizontal_end = horizontal_start + size
        vertical_start = j * size + j * margin
        vertical_end = vertical_start + size
        results[horizontal_start:horizontal_end, vertical_start:vertical_end] = filter_img.T

plt.figure(figsize=(14,8))
plt.imshow(results, cmap='gray')
plt.title(layer_name)
plt.tick_params(axis='both', which='both', bottom=False, left=False, labelleft=False, labelbottom=False)
plt.show()
```


```
In [22]: plt_filter_patterns(model, 'conv2d_2') # 128 x 128 (downward conv2d)
```




```
In [23]: plt_filter_patterns(model, 'conv2d_4') # 64 x 64 (downward conv2d)
```



```
In [24]: plt_filter_patterns(model, 'conv2d_6') # 32 x 32 (downward conv2d)
```



```
In [25]: plt_filter_patterns(model, 'conv2d_22') # 128 x 128 (upward conv2d)
```



Prediction & Down sampling

```
In [26]: def downsample(img):  
    ''' Downsample images from 128 x 128 to 100 x 100 '''  
    if size_img == size_img_target:  
        return img  
    return resize(img, (size_img, size_img), mode='constant', preserve_range=True)  
  
util_downsample = lambda im: np.array([downsample(x) for x in im])
```

```
In [27]: lb_pred_data = util_downsample(model.predict(data_img_up))
lb_pred_train = util_downsample(model.predict(im_train))
lb_orig_train = util_downsample(lb_train)
lb_pred_val = util_downsample(model.predict(im_val))
lb_orig_val = util_downsample(lb_val)
lb_pred_test = util_downsample(model.predict(test_img_up))
```

Find optimal threshold based on IoU

```
In [28]: def iou_metric(y_true_in, y_pred_in, print_table=False):
''' Function to calculate Intersection of Intersection of Union (IoU) '''
    labels = y_true_in
    y_pred = y_pred_in
    true_objects = 2
    pred_objects = 2
    intersection = np.histogram2d(labels.flatten(), y_pred.flatten(), bins=(true_objects, pred_objects))[0]

    # Compute areas (needed for finding the union between all objects)
    area_true = np.histogram(labels, bins = true_objects)[0]
    area_pred = np.histogram(y_pred, bins = pred_objects)[0]
    area_true = np.expand_dims(area_true, -1)
    area_pred = np.expand_dims(area_pred, 0)

    # Compute union
    union = area_true + area_pred - intersection

    # Exclude background from the analysis
    intersection = intersection[1:,1:]
    union = union[1:,1:]
    union[union == 0] = 1e-9

    # Compute the intersection over union
    iou = intersection / union

    # Loop over IoU thresholds
    prec = []
    if print_table:
        print('Thresh\tTP\tFP\tFN\tPrec.')
    for t in np.arange(0.5, 1.0, 0.05):
        tp, fp, fn = precision_at(t, iou)
        if (tp + fp + fn) > 0:
            p = tp / (tp + fp + fn)
        else:
            p = 0
        if print_table:
            print('{:1.3f}\t{}\t{}\t{}\t{:1.3f}'.format(t, tp, fp, fn, p))
        prec.append(p)

    if print_table:
        print('AP\t-\t-\t-\t{:1.3f}'.format(np.mean(prec)))
    return np.mean(prec)
```


```
In [29]: def precision_at(threshold, iou):
''' Calculate True Positive/False Positive/False Negative '''
    matches = iou > threshold
    true_positives = np.sum(matches, axis=1) == 1 # Correct objects
    false_positives = np.sum(matches, axis=0) == 0 # Missed objects
    false_negatives = np.sum(matches, axis=1) == 0 # Extra objects
    tp, fp, fn = np.sum(true_positives), np.sum(false_positives), np.sum(false_negatives)
    return tp, fp, fn
```

```
In [30]: def iou_metric_batch(y_true_in, y_pred_in):
''' Wrapper for IoU calculatoin '''
    batch_size = y_true_in.shape[0]
    metric = []
    for batch in range(batch_size):
        value = iou_metric(y_true_in[batch], y_pred_in[batch])
        metric.append(value)
    return np.mean(metric)
```

```
In [31]: def plot_threshold_IoU(thresholds, ious, threshold_best, iou_best):
''' Plot IoU vs threshold '''
fig,ax = plt.subplots(figsize=(7,4))
plt.plot(thresholds, ious)
plt.plot(threshold_best, iou_best, 'xr', label='Best threshold')
plt.xlabel('Threshold',fontSize=12)
plt.ylabel('IoU',fontSize=12)
plt.title('Threshold vs IoU ({:.3f}, {:.3f})'.format(threshold_best, iou_best),fontSize=12)
plt.legend(fontsize=10,loc='lower right')
ax.tick_params(labelsize=12)
plt.show()
```

```
In [32]: thresholds = np.linspace(0, 1, 50)
ious = np.array([iou_metric_batch(lb_orig_val, np.int32(lb_pred_val > x)) for x in thresholds])
threshold_best_index = np.argmax(ious)
iou_best = ious[threshold_best_index]
threshold_best = thresholds[threshold_best_index]
```

```
In [33]: plot_threshold_IoU(thresholds, ious, threshold_best, iou_best)
```



Show some results

```
In [34]: def id2str(ids_pred):
''' Convert index of a mask image into a series of numbers '''
output = []
conseq = 1
for i in range(len(ids_pred)):
    if i == 0:
        output.append(ids_pred[i])
    if i == len(ids_pred) - 1:
        output.append(conseq)
    elif i == len(ids_pred) - 1:
        if ids_pred[i] - ids_pred[i-1] == 1:
            conseq += 1
        else:
            conseq = 1
        output.append(conseq)
    else:
        if ids_pred[i] - ids_pred[i-1] == 1:
            conseq += 1
        else:
            output.append(conseq)
            output.append(ids_pred[i])
            conseq = 1
return output
```

```
In [35]: def ids_output(y_pred,threshold):
''' Convert a saltbody probability image into index '''
y_label = copy.deepcopy(y_pred)
y_label[y_label > threshold] = 1
y_label[y_label <= threshold] = 0
ids_pred = np.where(np.array([int(i) for i in y_label]) == 1)[0]+ 1
ids_pred = ids_pred.tolist()
output = id2str(ids_pred)
return str(output).strip('[]').replace(',','')
```



```
In [36]: def str2img(strnum,size_img):
''' Convert a series of numbers into a mask image '''
img_label = np.zeros(size_img*size_img)
if len(strnum) > 0:
ints = list(map(int,strnum.split(' ')))
i = 0
for i in range(int(len(ints)/2)):
idx_start = ints[2*i] - 1
num_seq = ints[2*i+1]
img_label[idx_start:idx_start+num_seq] = 1
return img_label
```

```
In [37]: def show_img(img_data,lb_pred,threshold,idx_img,training=True):
''' Function to display example images '''

# Prepare images
resaper = lambda x,y: np.reshape(x,[y,y])
img_seis = img_data[idx_img,:,:]
img_pred = resaper(lb_pred[idx_img,:],size_img)
img_cutoff = copy.deepcopy(img_pred)
img_cutoff[img_cutoff > threshold] = 1
img_cutoff[img_cutoff <= threshold] = 0


if training:
title_prefix = 'Training (Index: ' + str(idx_img) + ')'
idx_mask = train[train['id']==ids_train[idx_img]].rle_mask.values[0]
img_label = resaper(str2img(idx_mask,size_img),size_img)
else:
title_prefix = 'Test (Index: ' + str(idx_img) + ')'
idx_mask = ids_output(lb_pred[idx_img,:],threshold)
img_label = np.zeros_like(img_seis)

# Display images
fig, axes = plt.subplots(1, 4, figsize=(14, 5))
for i,ax in enumerate(axes.flat):
plt.axes(ax)
plt.imshow(img_seis.T,cmap='gray')
if i == 0:
plt.imshow(np.zeros_like(img_seis).T,alpha=0.5,clim=[0,1],cmap='Blues')
plt.title('Seismic')
elif i == 1:
plt.imshow(img_pred.T,alpha=0.5,clim=[0,1],cmap='Blues')
plt.title('Salt Probability')
elif i == 2:
plt.imshow(img_cutoff.T,alpha=0.5,clim=[0,1],cmap='Blues')
plt.title('Predicted Mask')
elif i == 3:
plt.imshow(img_label.T,alpha=0.5,clim=[0,1],cmap='Blues')
plt.title('True Mask')

plt.tick_params(axis='both',which='both',bottom=False,left=False,labelleft=False,labelbottom=False)


fig.subplots_adjust(right=0.925)
cbar_ax = fig.add_axes([0.95, 0.25, 0.01, 0.5])
clb = plt.colorbar(ticks=np.arange(0,1+0.001,.2),cax=cbar_ax)
text_ax = fig.add_axes()
plt.text(-50,-0.1,title_prefix,fontsize=14)
plt.show()
```

```
In [38]: show_img(data_img,lb_pred_data,threshold_best,idx_img=1500)
```



Prediction for test dataset

```
In [39]: lb_pred_test = lb_pred_test.reshape([np.shape(test_img_up)[0],size_img**2])  
show_img(test_img,lb_pred_test,threshold_best,idx_img=350,training=False)
```



Output for submission

```
In [40]: def output2csv(sample_submit,y_pred_test,ids_test,threshold):  
        ''' Output a csv file for submission '''  
        test_submit = sample_submit.copy()  
        num_output = len(test_submit)  
        for i in range(num_output):  
            id_tmp = np.where(test_submit.id.values == ids_test[i])[0][0]  
            test_submit['rle_mask'].iloc[id_tmp] = ids_output(y_pred_test[i,:],threshold)  
        test_submit.to_csv('./test.csv',index=False)
```

```
In [41]: output2csv(sample_submit,lb_pred_test,ids_test,threshold_best)
```